

Gitesh Chawda

Machine Learning Engineer

(657) 681-8005 | gitesh.ch.0912@gmail.com | linkedin.com/in/gitesh-ch | github.com/IMvision12 | imvision12.vercel.app

SUMMARY

Software and Machine Learning Engineer with hands-on experience building and deploying production-grade backend and frontend systems across cloud environments. Experienced in designing and scaling LLM pipelines, agentic AI systems, and full-stack AI products with real-world deployment. Strong background in deep learning, computer vision, and transformer-based architectures with contributions across the Hugging Face and Keras ecosystems.

TECHNICAL SKILLS

Languages: Python, C++, JavaScript, TypeScript, SQL

ML and Deep Learning: PyTorch, TensorFlow, Keras 3, JAX, Hugging Face Transformers, scikit-learn, Transformers, Flash Attention, Diffusion Models, Computer Vision

LLMs, Inference, and Fine-tuning: vLLM, TensorRT, Triton, CUDA, XLA, bitsandbytes, Quantization, LoRA, PEFT, SFT, GRPO, KV Cache Optimization, Model Serving

Agentic AI Frameworks: LangChain, LangGraph, LlamaIndex, CrewAI, AutoGen, OpenAI Agents SDK, MCP, A2A

MLOps and Cloud: Docker, Linux, Git, GitHub Actions, CI/CD, MLOps, LLMOps, AWS, GCP

Web and Data: React.js, Flask, GraphQL, Node.js, NumPy, Pandas, OpenCV, MySQL, MongoDB, Oracle

EXPERIENCE

KerasFormers, Author and Maintainer, Open Source (2024 - Present) [Github](#)

- Built a backend-agnostic, pip-installable Keras 3 library supporting 60+ architectures across vision, speech, LLM, and VLM tasks (SAM, DETR, DINOv3, Whisper, Llama, Qwen, DeepSeek, Mistral, Gemma, GLM), with single-codebase execution on TensorFlow, PyTorch, and JAX.
- Engineered Hugging Face-to-Keras checkpoint conversion with bit-close parity validation against reference implementations, and developed Keras 3 native weight-only quantization (INT8, INT4, FP8) using backend-agnostic operations, enabling up to 8x model size reduction for efficient inference deployment.

Open Source Contributor, Hugging Face Transformers, Keras, Keras.io, KerasCV (2022 - Present)

- Authored 42 merged PRs across Hugging Face Transformers, Keras, Keras.io, and KerasCV, improving model APIs, tokenizer design, type safety, and developer workflows.
- Added 10+ Keras 3 multi-backend APIs, layers, and utilities including PSNR, unravel_index, diagflat, RandomGrayscale, Equalization, and image augmentation workflows.
- Enhanced Hugging Face and Keras ecosystems through tokenizer modularity, type hints, ConvMixer models, YOLOv8 improvements, detection tests, segmentation augmentations, and official keras.io examples.

Fiserv, Software Developer, Chennai, India (July 2023 - August 2024)

- Developed and maintained a production web application with React.js, GraphQL, and AWS (AppSync, Lambda, API Gateway, Cognito, S3), supporting internal banking and payments workflows.
- Strengthened application security by implementing client-side data tampering prevention and request validation, while collaborating with AppSec teams to remediate vulnerabilities identified through internal security audits and penetration testing.
- Managed 20+ production deployments and critical hotfixes across development, staging, and production environments while maintaining reliable application availability.

Braynix AI, Deep Learning Intern, Bangalore, India (September 2021 - December 2021)

- Built deep learning pipelines for medical X-ray abnormality detection, improving validation accuracy by 20% over baseline CNN models across 5,000+ images.
- Reached 92% precision and 0.85+ IoU in medical image segmentation by implementing a U-Net with an EfficientNet-B4 encoder.
- Trained and benchmarked 3 object detection models (YOLOv5, SSD, RetinaNet) on 1,000+ labeled images at 0.78+ mAP.

PROJECTS

txtcode: TypeScript framework for orchestrating AI coding agents across 6 messaging platforms and 10 LLM backends. [Github](#)

Trinix: Triton-accelerated PyTorch optimization library with custom GPU kernels delivering up to 4.8x faster execution. [Github](#)

InferML: Local-first inference server for 143+ Hugging Face model families with an OpenAI-compatible API. [Github](#)

OpenRAG: Local-first RAG framework supporting vector and knowledge graph retrieval with local LLMs. [Github](#)

EDUCATION

California State University, Fullerton, Master of Science in Computer Science, California, USA (August 2024 - May 2026)

Pimpri Chinchwad College of Engineering, Bachelor of Engineering in Computer Engineering, Pune, India (August 2019 - July 2023)

PUBLICATIONS

EffiConvRes: An Efficient Convolutional Neural Network with Residual Connections and Depthwise Convolutions, 2023 7th International Conference on Computing, Communication, Control and Automation (ICCUBEA).

Survey and Comparison of Various Pre-trained CNN Architectures and CNN-Transformer Combinations, Smart Grid and Electric Vehicle Conference.